


Convolutional Neural Networks

+ Tensorflow Introduction




Imagenet (Fei-Fei Li, Stanford, 2006)


1990's [LeNet]: LeCune uses neural networks to read zip codes

2012 [AlexNet]: Blows away the competition in visual classification using ConvNets


2013-2017 [ZF Net, GoogLeNet, VGGNet, ResNet]:
ConvNets continue to evolve to human levels of accuracy



Fully connected vs. convolutional architecture



A typical 3 layer fully connected network



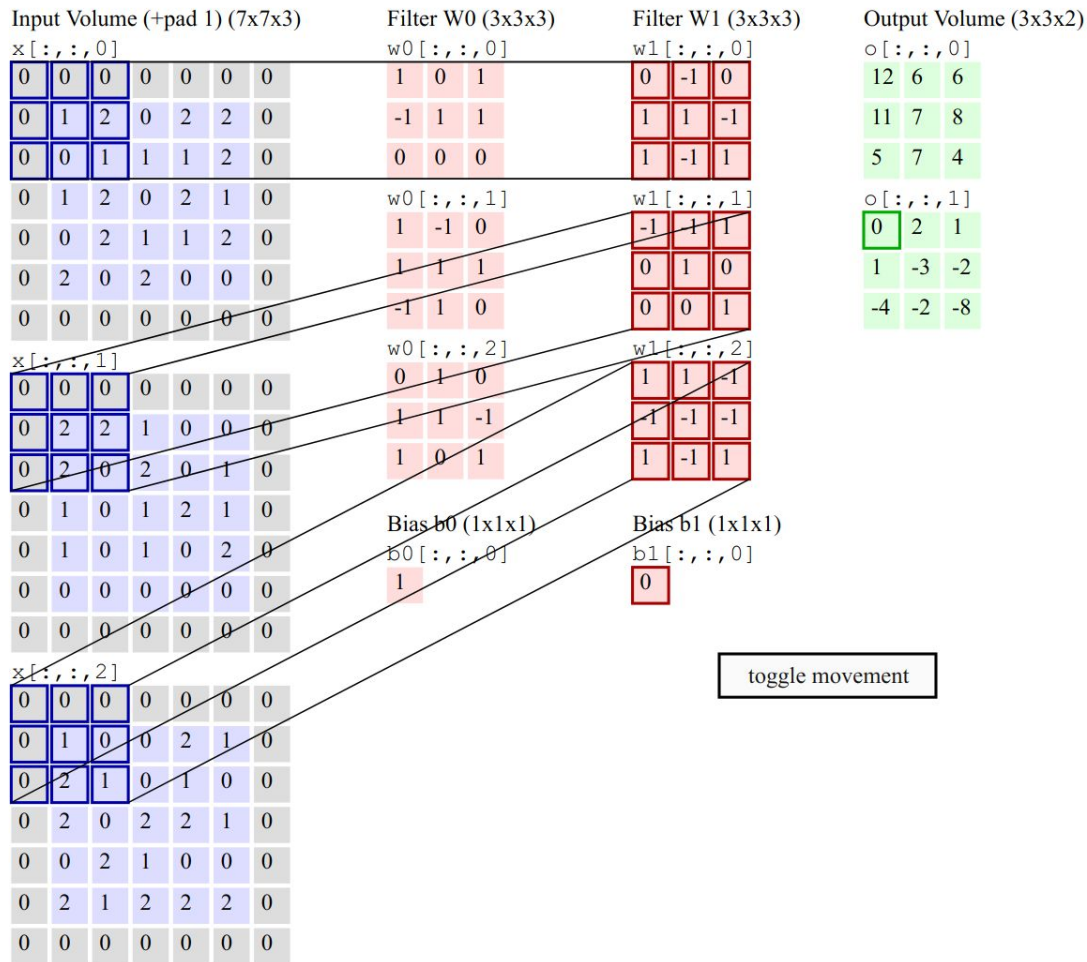
A convolutional network, the individual layers are a 3D tensor of shape:
`[kernel_height, kernel_width, features]`

Note: weights are not depicted

Convolutions


Stanford's online CS231n:

<http://CS231n.github.io/convolutional-networks/>




3D visualization of convolutional network

<http://scs.ryerson.ca/~aharley/vis/>

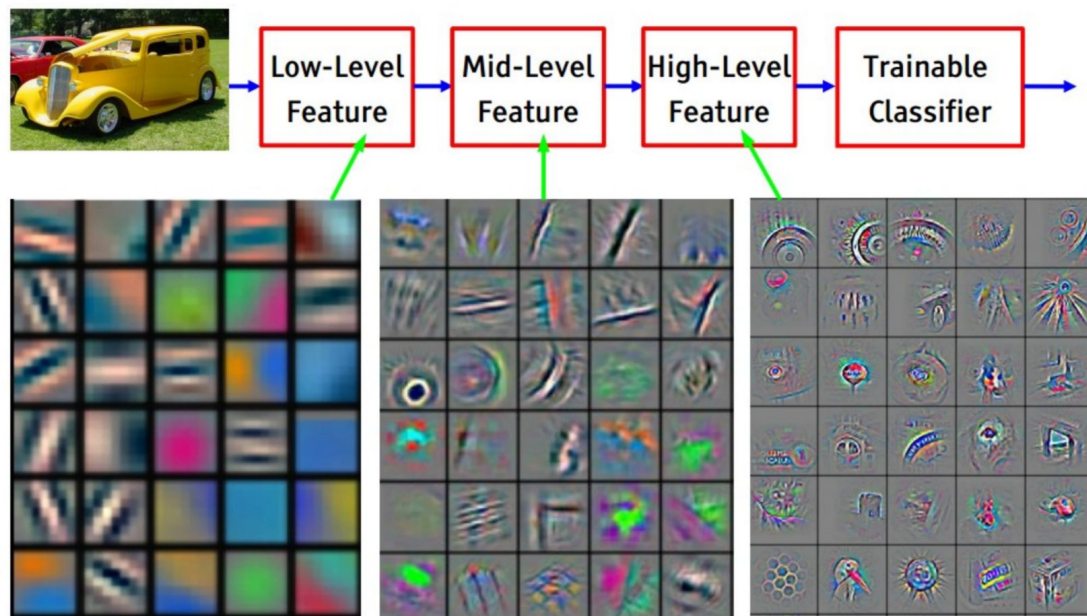


Reading CNN diagrams



Visualizing what was learned by the network

*[From recent Yann
LeCun slides]*




Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

Convolutional filters

The Convolution layer uses kernels to extract patterns from the image.

The kernels are learned during training

Pooling is a fancy method of downsampling



DeepViz

Open source visualization tool based on a number of papers circa 2014


“Deep Visualization Toolbox”
(4min)

“Visualizing and Understanding Deep Neural Networks by Matt Zeiler” (1hr)





0:49 (light -> dark filter)
1:00 (layer 1 filters)
2:33 (conv5 wrinkles filter)
3:05 (conv5 text filter)



VGGNet case study (ILSVRC 2014 runner up)



VGGNet case study (ILSVRC 2014 runner up)



INPUT:	[224x224x3]	memory:	224*224*3=150K	weights:	0
CONV3-64:	[224x224x64]	memory:	224*224*64=3.2M	weights:	(3*3*3)*64 = 1,728
CONV3-64:	[224x224x64]	memory:	224*224*64=3.2M	weights:	(3*3*64)*64 = 36,864
POOL2:	[112x112x64]	memory:	112*112*64=800K	weights:	0
CONV3-128:	[112x112x128]	memory:	112*112*128=1.6M	weights:	(3*3*64)*128 = 73,728
CONV3-128:	[112x112x128]	memory:	112*112*128=1.6M	weights:	(3*3*128)*128 = 147,456
POOL2:	[56x56x128]	memory:	56*56*128=400K	weights:	0
CONV3-256:	[56x56x256]	memory:	56*56*256=800K	weights:	(3*3*128)*256 = 294,912
CONV3-256:	[56x56x256]	memory:	56*56*256=800K	weights:	(3*3*256)*256 = 589,824
CONV3-256:	[56x56x256]	memory:	56*56*256=800K	weights:	(3*3*256)*256 = 589,824
POOL2:	[28x28x256]	memory:	28*28*256=200K	weights:	0
CONV3-512:	[28x28x512]	memory:	28*28*512=400K	weights:	(3*3*256)*512 = 1,179,648
CONV3-512:	[28x28x512]	memory:	28*28*512=400K	weights:	(3*3*512)*512 = 2,359,296
CONV3-512:	[28x28x512]	memory:	28*28*512=400K	weights:	(3*3*512)*512 = 2,359,296
POOL2:	[14x14x512]	memory:	14*14*512=100K	weights:	0
CONV3-512:	[14x14x512]	memory:	14*14*512=100K	weights:	(3*3*512)*512 = 2,359,296
CONV3-512:	[14x14x512]	memory:	14*14*512=100K	weights:	(3*3*512)*512 = 2,359,296
CONV3-512:	[14x14x512]	memory:	14*14*512=100K	weights:	(3*3*512)*512 = 2,359,296
POOL2:	[7x7x512]	memory:	7*7*512=25K	weights:	0
FC:	[1x1x4096]	memory:	4096	weights:	7*7*512*4096 = 102,760,448
FC:	[1x1x4096]	memory:	4096	weights:	4096*4096 = 16,777,216
FC:	[1x1x1000]	memory:	1000	weights:	4096*1000 = 4,096,000

TOTAL memory: 24M * 4 bytes ~ 93MB / image (only forward! ~*2 for bwd)


TOTAL params: 138M parameters

Question

Convolutional networks changed the world
(thank the Imagenet competition)

Why don't fully connected networks perform as well as ConvNets?

Fooling the network (optional)




Further Resources

Convolutional Neural Networks:

- neuralnetworksanddeeplearning.com
- Stanford's CS231n online deep learning course
- Deep Learning Book, MIT press (online & in print)
- Neural Network Playground: playground.tensorflow.org

Tensorflow

What is a “Tensor”



The most common error in tensorflow:

`ValueError: Shape must be rank 2 but is rank 3 for 'MatMul' (op: 'MatMul') with input shapes: [10,512,1], [512,2]`

Rank of a tensor	Math Entity	Example
0	Scalar	$x = 42$
1	Vector	$z = [10, 15, 20]$
2	Matrix	$a = \begin{bmatrix} 1 & 0 & 2 & 3 \\ 2 & 1 & 0 & 4 \\ 0 & 2 & 1 & 1 \end{bmatrix}$
3	3-Tensor (a cube of numbers)	A single image of shape [height, width, color_channels] example: [1080, 1920, 3]
4	4-Tensor (a set of cubes)	A batch of n images with shape [batch_size, height, width, channels] example: [10, 1080, 1920, 3]
n	n-dimensional Tensor	You get the idea...

Tensorflow is a math library

A general purpose math library created for gradient based operations

- Gradients are computed automatically
 - You can trivially ask for the gradient of any variable w.r.t. another
- Automatically handles optimized computation on GPU or CPU seamlessly
- Tensorflow is a symbolic computing library, not an imperative language
 - `if` and `while` loops don't function the same
- Provides a comprehensive set of operations used in constructing neural networks
- TF is NOT limited to neural networks!
- Tools for Visualization, Debugging, Profiling

Tensorflow constructs:

- Placeholders (Inputs)
- Tensors (Mutable variable)
- OPS (Operations)
- Constants

- Computation Graph data structure
- Session

Datatypes:

- `tf.float32` is standard
- Supports many data types similar to numpy

Basic workflow for Tensorflow

1 Build a computation Graph

```
import tensorflow as tf
```

```
a = tf.constant(2.0, tf.float32, name='a')  
b = tf.constant(3.0, tf.float32, name='b')
```

```
c = tf.multiply(a, b)
```

```
<tf.Tensor 'a:0' shape=() dtype=float32>
```

```
<tf.Tensor 'b:0' shape=() dtype=float32>
```

```
<tf.Tensor 'Mul_2:0' shape=() dtype=float32>
```

2 Run computations on the graph

```
sess = tf.Session()  
result = sess.run([b, c])
```

```
print(result[0])  
3.0
```

```
print(result[1])  
6.0
```

Passing parameters to tensorflow

```
import tensorflow as tf

# Build your graph
a = tf.placeholder(tf.float32, shape=(), name='a') # A scalar
b = tf.placeholder(tf.float32, shape=(2), name='b') # A vector
c = tf.multiply(a, b, name='c')

# Launch a session & evaluate tensor c
with tf.Session() as sess:
    params = {a: 10, b: [1, 2]} # a & b are tensor objects
    result = sess.run([c], feed_dict=params)
```

```
>>> print result
[array([ 10.,  20.], dtype=float32)]
```

Best practices note:

Structure your code with a `build_graph(...)` function which separates tensorflow graph operations from iterating through a dataset with calls to `sess.run(...)`

Demos

<https://github.com/aymericdamien/TensorFlow-Examples>

Tutorial index

0 - Prerequisite

- [Introduction to Machine Learning](#).
- [Introduction to MNIST Dataset](#).

1 - Introduction

- [Hello World](#) ([notebook](#)) ([code](#)). Very simple example to learn how to print "hello world" using TensorFlow.
- [Basic Operations](#) ([notebook](#)) ([code](#)). A simple example that cover TensorFlow basic operations.

2 - Basic Models

- [Linear Regression](#) ([notebook](#)) ([code](#)). Implement a Linear Regression with TensorFlow.
- [Logistic Regression](#) ([notebook](#)) ([code](#)). Implement a Logistic Regression with TensorFlow.
- [Nearest Neighbor](#) ([notebook](#)) ([code](#)). Implement Nearest Neighbor algorithm with TensorFlow.
- [K-Means](#) ([notebook](#)) ([code](#)). Build a K-Means classifier with TensorFlow.
- [Random Forest](#) ([notebook](#)) ([code](#)). Build a Random Forest classifier with TensorFlow.

3 - Neural Networks

Supervised

- [Simple Neural Network](#) ([notebook](#)) ([code](#)). Build a simple neural network (a.k.a Multi-layer Perceptron) to classify MNIST digits dataset. Raw TensorFlow implementation.
- [Simple Neural Network \(tf.layers/estimator api\)](#) ([notebook](#)) ([code](#)). Use TensorFlow 'layers' and 'estimator' API to build a simple neural network (a.k.a Multi-layer Perceptron) to classify MNIST digits dataset.
- [Convolutional Neural Network](#) ([notebook](#)) ([code](#)). Build a convolutional neural network to classify MNIST digits dataset. Raw TensorFlow implementation.
- [Convolutional Neural Network \(tf.layers/estimator api\)](#) ([notebook](#)) ([code](#)). Use TensorFlow 'layers' and 'estimator' API to build a convolutional neural network to classify MNIST digits dataset.
- [Recurrent Neural Network \(LSTM\)](#) ([notebook](#)) ([code](#)). Build a recurrent neural network (LSTM) to classify MNIST digits dataset.
- [Bi-directional Recurrent Neural Network \(LSTM\)](#) ([notebook](#)) ([code](#)). Build a bi-directional recurrent neural network (LSTM) to classify MNIST digits dataset.
- [Dynamic Recurrent Neural Network \(LSTM\)](#) ([notebook](#)) ([code](#)). Build a recurrent neural network (LSTM) that performs dynamic calculation to classify sequences of different length.

Unsupervised

- [Auto-Encoder](#) ([notebook](#)) ([code](#)). Build an auto-encoder to encode an image to a lower dimension and re-construct it.
- [Variational Auto-Encoder](#) ([notebook](#)) ([code](#)). Build a variational auto-encoder (VAE), to encode and generate images from noise.
- [GAN \(Generative Adversarial Networks\)](#) ([notebook](#)) ([code](#)). Build a Generative Adversarial Network (GAN) to generate images from noise.
- [DCGAN \(Deep Convolutional Generative Adversarial Networks\)](#) ([notebook](#)) ([code](#)). Build a Deep Convolutional Generative Adversarial Network (DCGAN) to generate images from noise.

4 - Utilities

- [Save and Restore a model](#) ([notebook](#)) ([code](#)). Save and Restore a model with TensorFlow.
- [Tensorboard - Graph and loss visualization](#) ([notebook](#)) ([code](#)). Use Tensorboard to visualize the computation Graph and plot the loss.
- [Tensorboard - Advanced visualization](#) ([notebook](#)) ([code](#)). Going deeper into Tensorboard; visualize the variables, gradients, and more...

5 - Data Management

- [Build an image dataset](#) ([notebook](#)) ([code](#)). Build your own images dataset with TensorFlow data queues, from image folders or a dataset file.
- [TensorFlow Dataset API](#) ([notebook](#)) ([code](#)). Introducing TensorFlow Dataset API for optimizing the input data pipeline.

6 - Multi GPU

- [Basic Operations on multi-GPU](#) ([notebook](#)) ([code](#)). A simple example to introduce multi-GPU in TensorFlow.
- [Train a Neural Network on multi-GPU](#) ([notebook](#)) ([code](#)). A clear and simple TensorFlow implementation to train a convolutional neural network on multiple GPUs.

Installing Tensorflow

Start with **Anaconda**, both python 2 and 3 are supported

All major python libraries are included: Numpy, Matplotlib, Jupyter Notebook, etc.

Then:

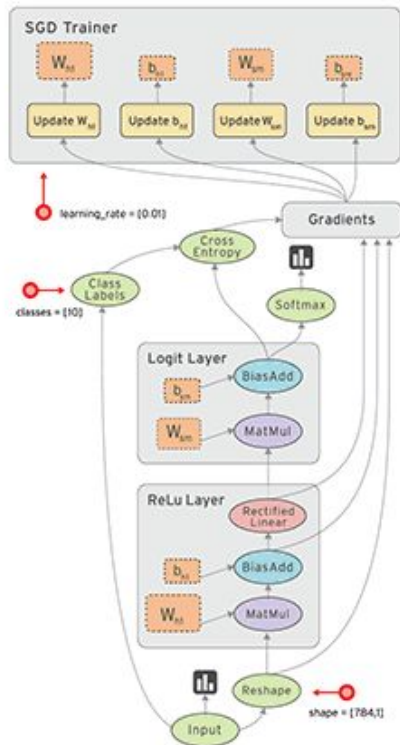
```
pip install tensorflow  
or  
pip install tensorflow-gpu
```

University GPU resources:

- **citrisdance.soe.ucsc.edu**
 - 2 older K20 GPUs and 32 cores storage is an issue
- **<https://patternlab.calit2.optiputer.net/>**
 - 2 fast M40 GPUs, 40 cores, shared with UCSD, 60+TB of storage
- More coming thanks to a recent NSF grant!

Automatic Differentiation


a.k.a.: How tensorflow and other tools do their magic



```
>>> import autograd.numpy as np # Thinly-wrapped numpy
>>> from autograd import grad   # The only autograd function you may ever need
>>>
>>> def tanh(x):                # Define a function
...     y = np.exp(-2.0 * x)
...     return (1.0 - y) / (1.0 + y)
...
>>> grad_tanh = grad(tanh)      # Obtain its gradient function
>>> grad_tanh(1.0)              # Evaluate the gradient at x = 1.0
0.41997434161402603
>>> (tanh(1.0001) - tanh(0.9999)) / 0.0002 # Compare to finite differences
0.41997434264973155
```

We can continue to differentiate as many times as we like, and use numpy's broadcasting of scalar-valued functions across many different input values:

```
>>> import matplotlib.pyplot as plt
>>> x = np.linspace(-7, 7, 200) # grad broadcasts across inputs
>>> plt.plot(x, tanh(x),
...         x, grad(tanh)(x), # first derivative
...         x, grad(grad(tanh))(x), # second derivative
...         x, grad(grad(grad(tanh)))(x), # third derivative
...         x, grad(grad(grad(grad(tanh)))(x), # fourth derivative
...         x, grad(grad(grad(grad(grad(tanh)))(x), # fifth derivative
...         x, grad(grad(grad(grad(grad(grad(tanh)))(x)) # sixth derivative
>>> plt.show()
```



<https://github.com/HIPS/autograd>